

# ECE 391 Exam 1, Fall 2010

Tuesday 28 September

Name:

- Be sure that your exam booklet has 11 pages.
- Write your name at the top of each page.
- This is a closed book exam.
- You are allowed one  $8.5 \times 11$ " sheet of notes.
- Absolutely no interaction between students is allowed.
- Show all of your work.
- Don't panic, and good luck!

*Name:* \_\_\_\_\_

2

Problem 1    25 points    \_\_\_\_\_

Problem 2    25 points    \_\_\_\_\_

Problem 3    25 points    \_\_\_\_\_

Problem 4    25 points    \_\_\_\_\_

Total        100 points    \_\_\_\_\_

Name: \_\_\_\_\_

3

**Problem 1 (25 points): Short Answers**

Please answer concisely. If you find yourself writing more than a sentence or two, your answer is probably wrong.

**Part A (5 points):** Recall the user-level test harness provided for your use with MP1. Describe one advantage and one disadvantage of developing and using such a testing strategy when writing new kernel code, relative to doing all testing of the new code directly in the kernel.

**Part B (5 points):** When an interrupt occurs in a generic cascade PIC configuration, both the master and slave don't know the correct vector number to send to the CPU. Explain why this is from each of the individual PIC's viewpoints and how this issue is addressed in the 8259A PIC design?

Name: \_\_\_\_\_

4

**Problem 1, continued:**

**Part C** (5 points): After a call returns from a C function, why should you remove its arguments off the stack AND not reuse those values?

**Part D** (5 points): Your friend suggests that mapping the 8259A interrupts into the 0x00 to 0x0F range of the Interrupt Descriptor Table can save time on translations between vector number and IRQ. Explain why such a mapping will not work as intended.

Name: \_\_\_\_\_

5

**Problem 1, continued:**

**Part E (5 points):** Several programs call the functions below in an arbitrary order. These programs share a single copy of the global variables `lock` and `rnum`. Can the call to `printf` ever print a number less than 1000? Explain.

```
spinlock_t      lock = SPIN_LOCK_UNLOCKED;
unsigned int     rnum = 0;

void generate()
{
    spin_lock (&lock);
    rnum = rand();    // Generate a random number from 0 to (2^32 - 1)
    spin_unlock (&lock);
}

void check()
{
    int cond = 0;

    spin_lock (&lock);
    if (rnum >= 1000)
        cond = 1;
    spin_unlock (&lock);

    if (1 == cond)
        printf("the number is %d\n", rnum)
}
```

Name: \_\_\_\_\_

6

## Problem 2 (25 points): Calling Conventions and the Stack

### Part A (10 points):

To improve the search speed of the `mpl_blink_struct` list, the linked list is to be converted into an array. You have been provided a skeleton `convert_to_array` function that takes in the current number of elements in the singly-linked list (`list_length`), allocates enough space for an array that will contain all of the current list nodes, fills in the array, and returns a pointer to the array.

```
void add_to_array(mpl_blink_struct *array, mpl_blink_struct *current,
                 int filled_array_elements);

void* mpl_malloc(int size);

mpl_blink_struct* convert_to_array(int list_length);
```

Complete `convert_to_array` by filling in the call to `add_to_array` following the appropriate calling convention. The code is on the following page. If you find yourself writing more than 15 lines of code, you have probably made a mistake.

`convert_to_array` will iterate through every element starting at `mpl_list_head` and call the `add_to_array` function. `add_to_array` takes:

- `array` - the pointer to the allocated array of `mpl_blink_structs`
- `current` - the `mpl_blink_struct` to add to array
- `filled_array_elements` - the current number of elements copied into array

Name: \_\_\_\_\_

7

## Problem 2, continued:

```
.long mp1_list_head
```

```
convert_to_array:
```

```
    pushl %ebp
    movl %esp,%ebp
    pushl %ebx
    pushl %esi
    pushl %edi
    movl 8(%ebp), %eax
    movl mp1_list_head, %edx
    imull $STRUCT_SIZE, %eax
    pushl %eax
    call mp1_malloc
    addl $4, %esp
    cmpl $0, %eax
    je MALLOC_FAIL
    movl $0, %ecx
```

```
CHECK_NEXT:
```

```
    cmpl $0, %edx
    je DONE_INSERTING
    /* Insert call to add_to_array here */
```

```
    incl %ecx
    movl %edi, %edx
    jmp CHECK_NEXT
```

```
DONE_INSERTING:
```

```
    popl %edi
    popl %esi
    popl %ebx
    leave
    ret
```

```
MALLOC_FAIL:
```

```
    pushl %eax
    call mp1_free
    addl $4, %esp
    movl $0, %eax
    jmp DONE_INSERTING
```

Name: \_\_\_\_\_


8

## Problem 2, continued:

### Part B (15 points):

The figures below are the state of the stack before the execution of the given code. Please fill in the state of the stack after execution and also indicate where `%ebp` and `%esp` are pointing to. Treat registers whose values you do not know as variables, otherwise fill in the actual number. We give you multiple figures for scratch work purposes but one must contain your final answer. **Circle the stack you want us to grade!**

**If you do not circle a stack, we will give you a zero on the problem.**



```
pushl $20
pushl %ebx
pushl %ecx
call foo
...
foo:
pushl %ebp
movl %esp, %ebp
pushl %ebx
pushl %esi
pushl %edi
addl $-8, %esp
movl 8(%ebp), %edi
movl 12(%ebp), %esi
xorl %ebx, %ebx
movl %ebx, 4(%esp)
addl $4, %esi
imull %esi, %esi
addl $-4, %esp
movl %esi, (%esp)
```



Name: \_\_\_\_\_

9

### Problem 3 (25 points): Synchronization

As part of your first job, you are charged with implementing a more fair lock construct. Specifically, you must implement a bakery lock (also called a ticket lock) based on the Linux semaphore.

A bakery lock is analogous to the physical mechanism used at many busy bakeries. When a customer enters the bakery, he/she takes the next paper ticket from the dispenser. Each ticket is marked with a unique number, and consecutive tickets have consecutive numbers (1, 2, 3, etc). A display indicates the ticket number currently being served. After taking a ticket, a customer watches the display for his/her number. When the display shows his/her number, the customer can order (that is, he/she owns the lock). After the customer orders (that is, when realizing the lock), the number displayed is incremented.

Using the following data structure and the Linux semaphore API (provided on the last page of the exam), implement the `bakery_init`, `bakery_lock`, and `bakery_unlock` functions and answer the subsequent questions.

**Note: Multiple people should be able to take a ticket while other people are waiting to be served**

```
struct bakery_lock_t {
    struct semaphore* sem;
    unsigned int dispenser;
    volatile unsigned int display;
};
typedef struct bakery_lock_t bakery_lock_t;
```

**Part A (4 points):** Implement `bakery_init` function below such that the display is initially set to 1. Be sure that the other fields are set properly to work with your answer to **Part B**.

```
void bakery_init (bakery_lock_t* b)
{
```

```
}
```

**Part B (7 points):** Implement `bakery_lock` function below. Return only after the bakery lock is owned by the caller.

```
void bakery_lock (bakery_lock_t* b)
{
```

```
}
```

Name: \_\_\_\_\_

10

**Problem 3, continued:**

**Part C (4 points):** Implement `bakery_unlock` function below.

```
void bakery_unlock (bakery_lock_t* b)
{
```

```
}
```

**Part D (5 points):** After you implement your lock, a friend in the company confides in you that they are nervous about the correctness of your implementation and would like to add a second synchronization variable to the bakery lock. Specifically, the friend suggests adding a reader-writer semaphore to the `bakery_lock_t` structure, acquiring a read lock in `bakery_lock` while waiting for the `display` field to show the caller's number, and acquiring a write lock in `bakery_unlock` before changing the display. Explain why you refuse to adopt your friend's suggestion.

**Part E (5 points):** Based on the design of the bakery lock in this problem, is it safe for a program that holds a bakery lock to acquire a Linux semaphore? Explain briefly why acquiring a semaphore (as opposed to a spin lock) might be a concern and why that particular problem will or will not show up with bakery lock.

Name: \_\_\_\_\_

11

#### Problem 4 (25 points): x86 Assembly

Your math assignment requires you to perform numerous dot products on sets of two vectors. Being a computer engineer, you decide to flex your knowledge of x86 assembly to write a dot product calculator and thus make your math assignment a joke. Write the function `dot_product` that iterates through two equal length singly-linked lists that represent the two vectors, and returns their dot product. The heads to the linked lists are `vect1_head` and `vect2_head`, which are global variables. The linked list structure is defined below:

```
struct vect_node_t {  
    vect_node_t* next;  
    int value;  
};
```

How to perform the dot product: given two vectors **A** and **B** where **A**={1,2,3} and **B**={4,5,6}, their dot product is  $(1 * 4) + (2 * 5) + (3 * 6) = 32$ . If you have any questions on the semantics of the dot product, please ask a TA.

Write your x86 assembly code on the next page. You may use the space provided on the current page to write the code in C, but your C code will not be graded. **Assume equal length vectors and the result will not overflow a 32 bit integer.**

*Name:* \_\_\_\_\_

12

**Problem 4, continued:**

```
.global vect1_head  
.global vect2_head  
  
dot_product:
```

*Name:* \_\_\_\_\_

13

(scratch paper)

Name: \_\_\_\_\_

14

**You may tear off this page to use as a reference**

## Synchronization API reference

<code>spinlock_t lock;</code>	Declare an uninitialized spinlock
<code>spinlock_t lock1 = SPIN_LOCK_UNLOCKED;</code> <code>spinlock_t lock2 = SPIN_LOCK_LOCKED;</code>	Declare a spinlock and initialize it
<code>void spin_lock_init(spinlock_t* lock);</code>	Initialize a dynamically-allocated spin lock (set to unlocked)
<code>void spin_lock(spinlock_t *lock);</code>	Obtain a spin lock; waits until available
<code>void spin_unlock(spinlock_t *lock);</code>	Release a spin lock
<code>void spin_lock_irqsave(spinlock_t *lock,                         unsigned long&amp; flags);</code>	Save processor status in <code>flags</code> , mask interrupts and obtain spin lock (note: <code>flags</code> passed by name (macro))
<code>void spin_lock_irqrestore(spinlock_t *lock,                           unsigned long flags);</code>	Release a spin lock, then set processor status to <code>flags</code>
<code>struct semaphore sem;</code>	Declare an uninitialized semaphore
<code>static DECLARE_SEMAPHORE_GENERIC (sem, val);</code>	Allocate statically and initialize to <code>val</code>
<code>DECLARE_MUTEX (mutex);</code>	Allocate on stack and initialize to one
<code>DECLARE_MUTEX_LOCKED (mutex);</code>	Allocate on stack and initialize to zero
<code>void sema_init(struct semaphore *sem, int val);</code>	Initialize a dynamically allocated semaphore to <code>val</code>
<code>void init_MUTEX(struct semaphore *sem);</code>	Initialize a dynamically allocated semaphore to one.
<code>void init_MUTEX_LOCKED(struct semaphore *sem);</code>	Initialize a dynamically allocated semaphore to zero.
<code>void down(struct semaphore *sem);</code>	Wait until semaphore is available and decrement (P)
<code>void up(struct semaphore *sem);</code>	Increment the semaphore

You may tear off this page to use as a reference

## x86 reference

8-bit			
32-bit	16-bit	high	low
EAX	AX	AH	AL
EBX	BX	BH	BL
ECX	CX	CH	CL
EDX	DX	DH	DL
ESI	SI		
EDI	DI		
EBP	BP		
ESP	SP		

← AX →

3116 158 70

AH

AL

← EAX →

jb	below	CF is set
jbe	below or equal	CF or ZF is set
je	equal	ZF is set
jle	less or equal	SF ≠ OF (SF ≠ OF) or ZF is set
jo	overflow	OF is set
jp	parity	PF is set (even parity)
js	sign	SF is set (negative)

movb (%ebp),%al# AL ← M[EBP]

movb -4(%esp),%al# AL ← M[ESP - 4]

movb (%ebx,%edx),%al# AL ← M[EBX + EDX]

movb 13(%ecx,%ebp),%al# AL ← M[ECX + EBP + 13]

movb (,%ecx,4),%al# AL ← M[ECX \* 4]

movb -6(,%edx,2),%al# AL ← M[EDX \* 2 - 6]

movb (%esi,%eax,2),%al# AL ← M[ESI + EAX \* 2]

movb 24(%eax,%esi,8),%al# AL ← M[EAX + ESI \* 8 + 24]

movb 100,%al# AL ← M[100]

movb label,%al# AL ← M[label]

movb label+10,%al# AL ← M[label+10]

movb 10(label),%al# NOT LEGAL!

movb label(%eax),%al# AL ← M[EAX + label]

movb 7\*6+label(%edx),%al# AL ← M[EDX + label + 42]

movw \$label,%eax# EAX ← label

movw \$label+10,%eax# EAX ← label+10

movw \$label(%eax),%eax# NOT LEGAL!

call printf# (push EIP), EIP ← printf

call \*%eax# (push EIP), EIP ← EAX

call \*(%eax)# (push EIP), EIP ← M[EAX]

call \*fptr# (push EIP), EIP ← M[fptr]

call \*10(%eax,%edx,2)# (push EIP), EIP ←

# M[EAX + EDX\*2 + 10]

Conditional branch sense is inverted by inserting an “N” after initial “J,” e.g., JNB. Preferred forms in table below are those used by debugger in disassembly. Table use: after a comparison such as  
cmp %ebx,%esi # set flags based on (ESI - EBX)  
choose the operator to place between ESI and EBX, based on the data type. For example, if ESI and EBX hold unsigned values, and the branch should be taken if ESI ≤ EBX, use either JBE or JNA. For branches other than JE/JNE based on instructions other than CMP, check the branch conditions above instead.

preferred form	jnz	jnae	jna	jz	jnb	jnbe	unsigned comparisons
	jne	jb	jbe	je	jae	ja	
	≠	<	≤	=	≥	>	
preferred form	jne	jl	jle	je	jge	jg	signed comparisons
	jnz	jnge	jng	jz	jnl	jnle	