

- reader/writer spin locks (rwlock_t)
 - extension of spin locks
 - use for short critical sections
 - ok to use in interrupt handlers
 - admit writer starvation (you must ensure that this not happen)
- reader/writer semaphores (struct rw_semaphore)
 - extension of semaphores
 - use only in system calls
 - do not admit writer starvation (new readers wait if writer is waiting)

type of code entering critical section	critical section shares data with	for mutual exclusion, use
system calls only	other system calls	up down
	interrupt handlers	spin_lock_irq spin_unlock_irq
both system calls and interrupt handlers	both system calls and interrupt handlers	spin_lock_irqsave spin_unlock_irqrestore
interrupt handlers only	system calls	spin_lock spin_unlock
	higher priority interrupt handlers	spin_lock_irqsave spin_unlock_irqrestore

	mutual exclusion	reader/writer	
data shared by interrupt handlers	spin_lock_irqsave	read_lock_irqsave	write_lock_irqsave
	spin_unlock_irqrestore	read_unlock_irqrestore	write_unlock_irqrestore
data shared only by system calls	down	down_read	down_write
	up	up_read	up_write

Hardware interrupts require some small amount of work to service a device, but may require much more work to handle any data delivered by the device or modified as a result of the interrupt

Software interrupts can interrupt programs, but can be interrupted by device // allow slower devices to get attention from processor in a timely manner w/o requiring processor to poll devices periodically

port(A=?)	info contained in Initialization Control Word
ICW1 0	start init, edge-triggered inputs, cascade mode, 4 ICWs
ICW2 1	high bits of vector #
ICW3 1	master: bit vector of slaves; slave: input pin on master
ICW4 1	ISA=x86, normal/auto EOI

INTA – strobes pin after receiving INT signal // PIC sees & writes interrupt vector to D
RD – high, processor will read data (vector #) from PIC
WR – After the interrupt handler finishes, tells the PIC it is done by writing to A & D // processor will write data (comm and, EOI) to PIC // done via OUT

■ PIC removes interrupt from the list of in-service interrupt and returns to waiting

SP – master(1) or slave(0) || **CS** – ADDR[1:15] // remove LSB

CAS -- identify a particular slave PIC to write to the data bus // output pins for master & input pins for slave

EOI – If interrupt handler fails to send EOI, PIC keeps masking all interrupts of equal/lower priority indefinitely

Ack – start of interrupt handling acknowledge receipt of interrupt // masks the interrupt on the PIC, then sends the EOI signal


End – called to end interrupt // enables interrupt (unmasks it) on PIC || **disable and enable** allow nested disabling and re-enabling of active interrupts

Spinlocks – ensures no interrupt & suitable for quick processes like timing // poll in a tight loop, doing nothing else until the resource is acquired

Semaphore – long critical section = longer waiting time; can put waiting processes to sleep & free processor to do other processes first; sys call; Up/Down

Memory mapped to two ports

- o Command port (e.g. 0x20)
- o Data port (MUST be Command Port + 1)



CPU - PIC Signals

- o INTR - Activated by the PIC upon interrupt
- o INTA - Pulsed by the CPU whenever
- o D - Bidirectional communication between CPU and PIC
 - Used in programming the PIC, sending EOI, etc.

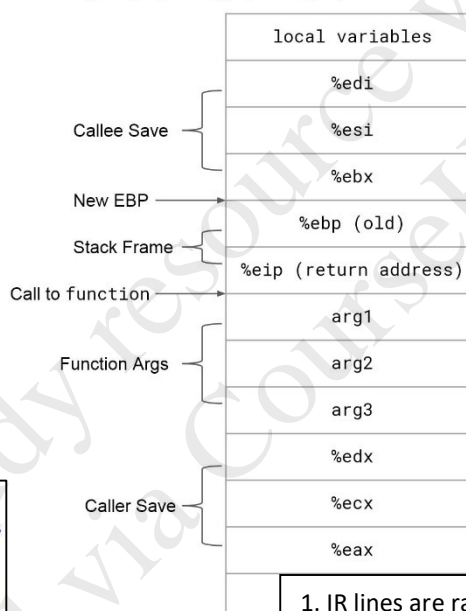
PIC Specific Signals

- o A - Distinguishes Command/Data port on PIC
 - Can be directly mapped to ADDR[0] (why?)
- o CS - Determines whether the given PIC should be active
 - Checks if ADDR[31:1] == PORT[31:1]

1. Lock & save flags
2. Mask interrupts to PIC so you don't get disturbed while initializing
3. Initialize PIC
4. Unmask interrupts
5. Unlock & restore flags

it's not good to take too much time in a hardware interrupt handler - other interrupts may need to do things too! That's what software interrupts are for - Software interrupts operate at priority level between regular programs and hardware interrupts, so hardware interrupts can generate a software interrupt to handle more time-intensive tasks, allowing other hardware interrupts to interrupt the software interrupts

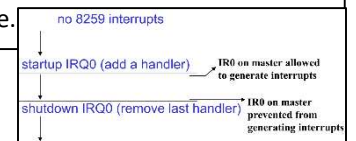
function(arg1, arg2, arg3):




Deadlock – waiting thread is still holding onto another resource that the first needs before it can finish // locks & unlocks are in reverse order
Livelock – thread tries to get lock, lets it go, then tries again // seems like it's working but it's stuck
Starvation – if many readers try to come in and writers can't write

8259A PIC

1. IR lines are raised that set corresponding IRR bits
2. PIC resolves priority and sends INT signal to CPU
3. The CPU acknowledges with INTA pulse.
4. When INTA signal received from CPU, highest priority ISR bit is set & corresponding IRR bit is reset. PIC does not drive data bus during this period.
5. The CPU will send second INTA pulse. PIC releases pointer to data bus from where it is read by the CPU.
6. ISR bit remains set until EOI command is issued at the end of interrupt subroutine.



```
int binary_search(int key, int* array, int size);
```



0. save caller-saved registers (if desired)

1. push arguments onto stack
2. make the call
3. pop arguments off the stack
4. restore caller-saved registers

-- After call, must restore stack ptr (ie pushl \$16, call, addl \$4, %esp) | store args from RTL

Volatile - value of the variable may change at any time w/o any action being taken by the code the compiler finds nearby
-- ESP instead of EBP: won't mess up base ptr, don't need to save old ebp; but no guarantee ESP is in right spot b/c it changes over time with stack

call vs jump:

- o jump → jmp LABEL
- o call → pushl %eip; jmp LABEL

enter - pushl %ebp; movl %esp, %ebp

- o "creates" the stack frame

leave - movl %ebp, %esp; popl %ebp

- o "tears down" the stack frame

ret - popl %eip

```
/* call the function */ # the call site...
value = afunc(10, 20);    pushl $20          # push second argument
                          pushl $10          # push first argument
                          call afunc         # call the function
                          addl $8,%esp       # pop the arguments
                          movl %eax,-4(%ebp) # store result in 'value'

int afunc(int a, int b)
{
    pushl %ebp           # save old frame pointer
    movl %esp,%ebp       # point to new frame
    subl $4,%esp         # make room for 'result'
    movl 8(%ebp),%eax     # put 'a' into EAX
    imull 12(%ebp),%eax   # multiply EAX by 'b'
    incl %eax            # add one
    movl %eax,-4(%ebp)   # store into 'result'
    movl -4(%ebp),%eax   # return 'result' in EAX
    leave               # restore frame pointer
    ret
}
```


Stack operations: The x86 ISA supports a stack abstraction directly rather than as a software convention. The PUSH and POP instructions provide the necessary functionality. As shown to the right, the stack convention used is that ESP contains the address of the element on top of the stack. The stack grows downward in addresses, like that of the LC-3, so

```
pushl %eax # M[ESP - 4] ← EAX, ESP ← ESP - 4
```

is equivalent to

```
movl %eax,-4(%esp) # M[ESP - 4] ← EAX
subl $4,%esp      # ESP ← ESP - 4
```

Other than a POP into the EFLAGS register, PUSH and POP do not affect the flags.



spin_lock:

```
movl 4(%esp), %eax
loop:
    movl $1,%ecx
    xchgl %ecx, (%eax)
    cmpl $1,%ecx
    je loop
    ret
```

spin_unlock:

```
movl 4(%esp), %eax
movl $0, (%eax)
ret
```

Initially, PIC interrupts masked out using mask
Startup -- 8259A tells the appropriate PIC to allow the interrupt line to generate interrupts; unmask the interrupt on PIC
Shutdown -- called when the last handler is removed from an interrupt & for 8259A, tells the appropriate PIC to mask interrupt

type	generated by	example	asynchronous	unexpected
interrupt	external device	packet arrived at network card	yes	yes
exception	invalid opcode or operand	divide by zero	no	yes
system call/trap	deliberate, via INT instruction	print character to console	no	no

What does the "P" mean on the "outb" macros?

- add PAUSE instruction after OUTB; "REP NOP" prior to P4
- delay needed for old devices that cannot handle processor's output rate

Critical section spans the whole function; why?

- avoid other 8259A interactions during initialization sequence
- (device protocol requires that four words be sent in order)

Why use **irqsave** for critical section?

- this code called from other interrupt initialization routines
- which may or may not have cleared IF on processor

- software-generated (soft) interrupt
- (similarly, but later, signals—user-level soft interrupts)
- runs at priority between program and hard interrupts
- usually generated

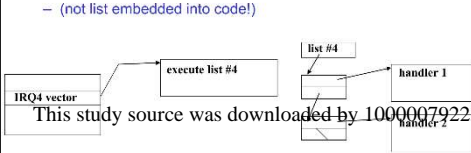
- by hard interrupt handlers
- to do work not involving device

Linux version is called **tasklets**

```
asm volatile (" # local_irq_save macro implementation
    pushfl # save EFLAGS to stack
    popl %0 # pop EFLAGS into output 0
    cli # mask interrupts
    : "=g" (flags) /* output 0 is a general-purpose register
    : /* which should then be stored in flags */
    : /* no inputs
    : /* no outputs
    : "g" (flags) /* input 0 is a general-purpose register
    : /* which should hold the value in flags */
    : "memory", "cc" /* see text
);

asm volatile (" # local_irq_restore macro implementation
    pushl %0 # save input 0 to stack
    popfl # pop input 0 into EFLAGS
    : /* no outputs
    : "g" (flags) /* input 0 is a general-purpose register
    : /* which should hold the value in flags */
    : "memory", "cc" /* see text
);
```

- interrupt chaining with linked list data structure
- (not list embedded into code!)



hw_irq_controller structure/struct
irq_chip -- each vector # associated with table; used to interact with an appropriate PIC

human-readable name

startup function

shutdown function

enable function


disable function

mask function

mask_ack function

unmask function

Kernel test -- safer & won't corrupt kernel-not using actual mem space that program will be running in
-invalid RW silenced in user level/will crash in kernel



on 8259

- first **disable_irq** calls jump table disable, which masks interrupt on PIC
- last **enable_irq** calls jump table enable, which unmasks interrupt on PIC

independent I/O -- use special instructions

and a separate I/O port address space

memory-mapped I/O -- use loads/stores

and dedicate part of the memory address space to I/O

Interrupt Chaining -- linked list of interrupt handlers associated w/ interrupt line

- for >1 device • must query devices to see if they raised interrupt • not always possible
- for 1 device • must avoid stealing data/confusing device (ie by sending two characters to serial port; in response to interrupt declaring port ready for one char; reading mouse location twice; if device protocol specifies reading once per interrupt)

0x00	division error
...	...
0x02	NMI (non-maskable interrupt)
0x03	breakpoint (used by KGDB)
0x04	overflow
...	...
0x0B	segment not present
0x0C	stack segment fault
0x0D	general protection fault
0x0E	page fault
...	...
0x20	IRQ0 -- timer chip
0x21	IRQ1 -- keyboard
0x22	IRQ2 -- (cascade to slave)
0x23	IRQ3
0x24	IRQ4 -- serial port (KGDB)
0x25	IRQ5
0x26	IRQ6
0x27	IRQ7
0x28	IRQ8 -- real time clock
0x29	IRQ9
0x2A	IRQ10
0x2B	IRQ11 -- eth0 (network)
0x2C	IRQ12 -- PS/2 mouse
0x2D	IRQ13
0x2E	IRQ14 -- ide0 (hard drive)
0x2F	IRQ15
0x30-0x7F	APIC vectors available to device drivers
0x80	system call vector (INT 0x80)
0x81-0xEE	more APIC vectors available to device drivers
0xEF	local APIC timer
0xF0-0xFF	symmetric multiprocessor (SMP) communication vectors

On full SMP systems, where multiple threads could be active in the kernel at the same time and interrupts could be delivered to pretty much any CPU, it's no longer enough to only disable interrupts on single processor, or only grab a single lock. Both are required: disabling interrupts protects from IRQ handler on the same CPU, holding a lock protects from other threads entering the same critical sections on different CPU. This is exactly why **spin_lock_irqsave()** and **spin_unlock_irqrestore()** were invented.