

# ECE391- Computer System Engineering

## Lecture 9

### Programable Interrupt Controller

## Office of the Provost

Dear Colleagues,

We have shared several updates about Fall 2021 in the past several weeks as we monitor the progress with managing the pandemic and respond to new science-based COVID-19 guidance from the CDC, IDPH, CUPHD and our own SHIELD team. Below is a condensed version of what we need to know as we prepare for in-person classes in a few weeks.

### Face Coverings

- Everyone (faculty, staff, students, visitors) is required to wear a face covering in university facilities. [Read more about face coverings and face shields here.](#)
- If a student is not wearing a face covering in your class, ask them to put one on. If they refuse, dismiss the class and report the student to the Office for Student Conflict Resolution for further discipline by [filling out this form](#). Call UIPD, 217-333-1216, only if an individual becomes belligerent, disruptive and threatening.

# Announcements

- PS2 Posted - Committed to the master (main) branch on GitLab by 5:59PM on 9/21
- MP2 Posted - All checkpoints should be committed to the master(main) branch on GitLab by:
  - Checkpoint 1: 5:59PM on 10/5
  - Final Checkpoint: 5:59PM on 10/12

# ECE391 EXAM 1

- EXAM I – Wednesday, September 29th, 7:00pm-9:00 PM
  - Location: ECEB 1002
- Online details to follow, but expect it to be proctored and at same time.
- NO Lecture on Tuesday, September 28
- Review Session
  - ?

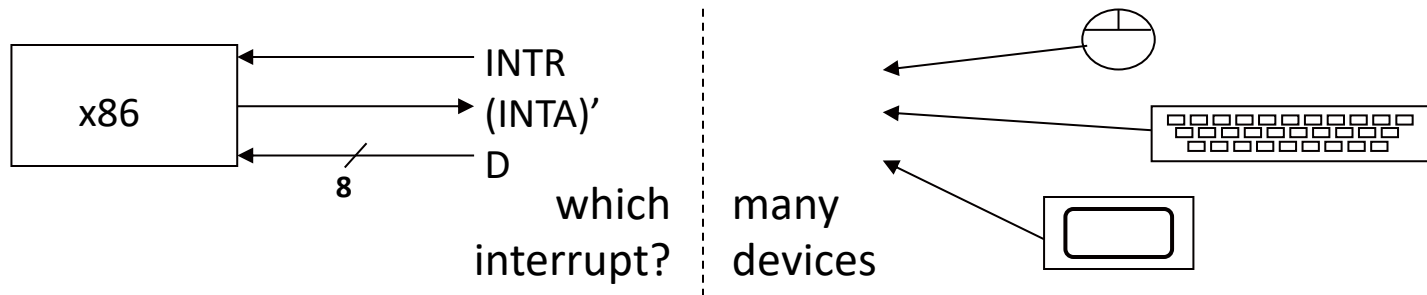
# Topics covered by EXAM 1

- Material covered in lectures (Lecture1 – Lecture10)
  - x86 Assembly
  - C (e.g., Calling Convention)
  - Synchronization
  - Interrupt control (using PIC)
- Material covered in discussions
- MP1

# Lecture Topics

- Programmable interrupt controller (PIC)
  - motivation & design
  - hardware for x86
  - Linux abstraction of PIC

# PIC Motivation and Design



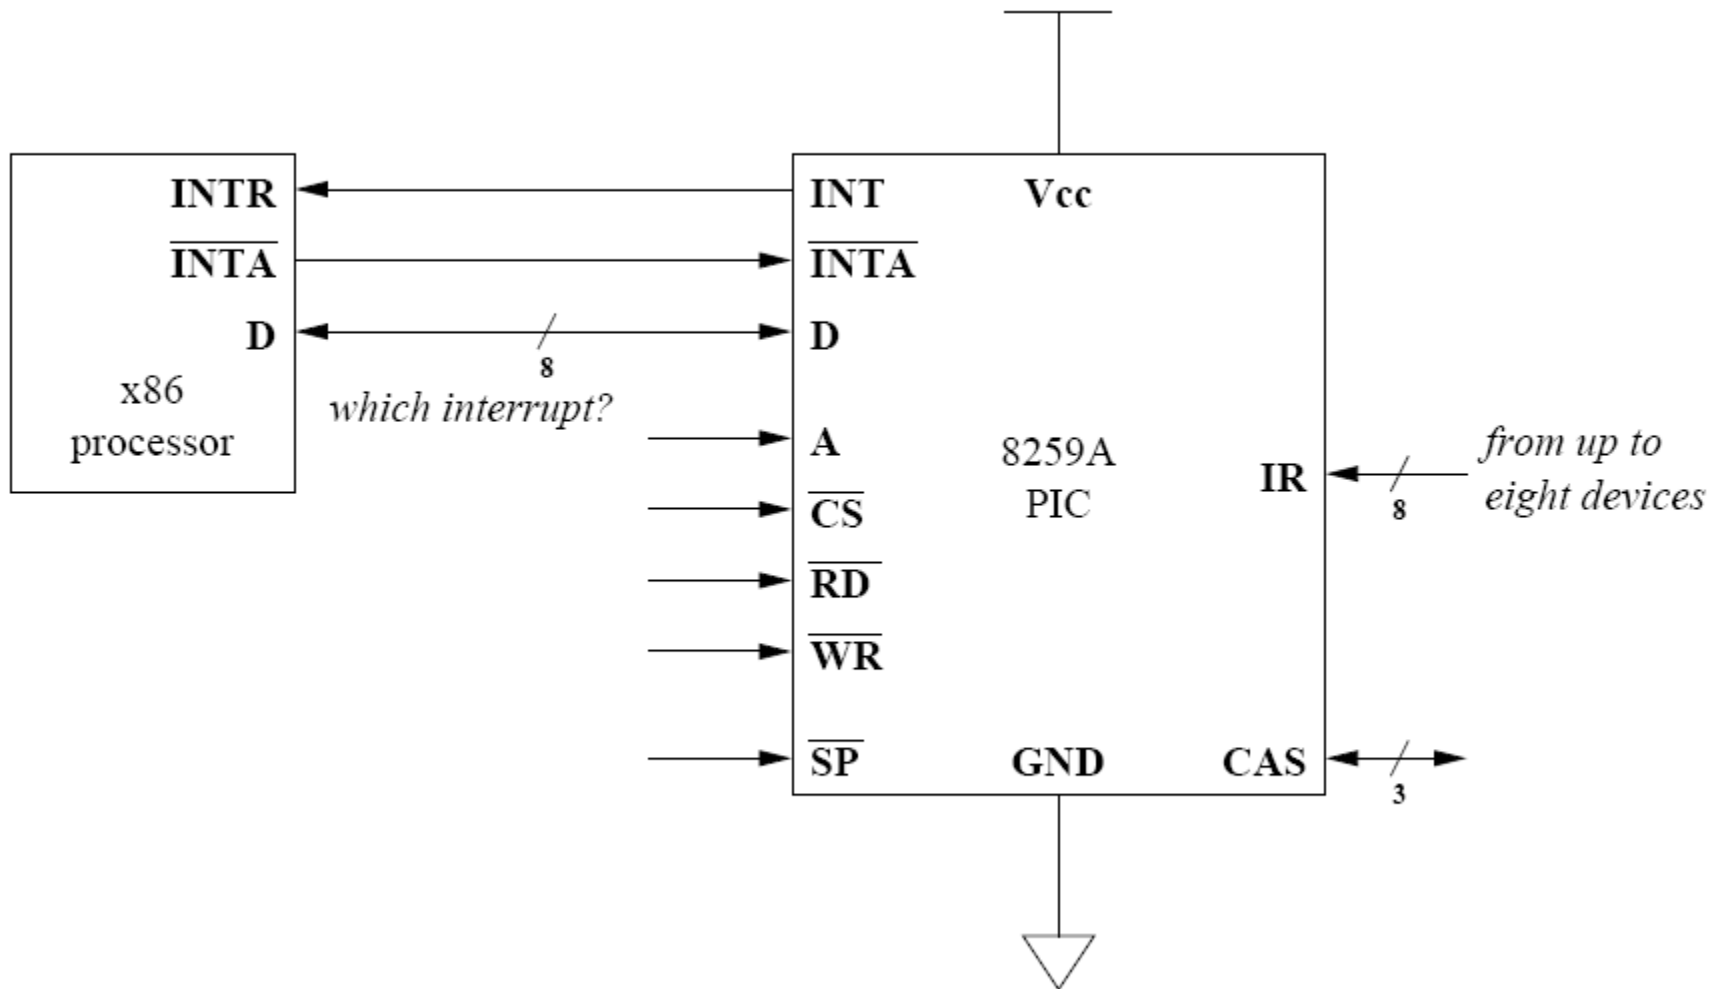
- How do we connect devices to the processor's interrupt input?
- An OR gate? why not?

# PIC Motivation and Design (cont.)



- who writes the vector # ?
  - possible to build arbiter, but...
  - what if more than one raised interrupt?
- extra work for processor to query all devices
  - must execute interrupt code for device that raised interrupt
  - could have been more than one device
  - no way to tell with OR gate
- not all devices support query
  - many devices too simplistic to support query
  - and operations (e.g., reading from port) may not be idempotent
- might be nice to have concept of priority and preemption, i.e., interrupting an interrupt handler



# 8259A Programmable Interrupt Controller (PIC)



# Intel 8259



[images: Wikipedia users German, Nixdorf]

# Logical Model of PIC Behavior

- Watch for interrupt signals
  - up to eight devices
  - one interrupt line each

# Logical Model of PIC Behavior

- Track which devices/input lines are currently in service by processor
  - i.e., processor is executing interrupt handler for that interrupt
  - using internal state

# Logical Model of PIC Behavior (cont.)

- When a device raises an interrupt
  - If priority is higher than those of in-service interrupts
    - report the highest-priority raised to the processor
    - mark that device as being in service
  - else
    - do nothing

# Logical Model of PIC Behavior (cont.)

- When processor reports EOI (end of interrupt) for some interrupt
  - remove the interrupt from the in-service mask
  - check for raised interrupt lines that should be reported to processor

# Logical Model of PIC Behavior (cont.)

- Protocol for reporting interrupts
  - PIC raises INTR
  - processor strobes INTA' (active low) repeatedly
    - creates cycles for PIC to write vector to data bus
    - (must follow spec timing! PIC is not infinitely fast!)
  - processor sends EOI with specific combinations of A & D inputs (A is from address bus, D is from data bus)


# Logical Model of PIC Behavior (cont.)

- What about A, CS', RD', and WR' ?
  - A = address match for ports
  - CS' = chip select (does processor want PIC to read/write?)
  - RD' and WR' defined from processor's point of view
    - RD' = processor will read data (vector #) from PIC
    - WR' = processor will write data (command, EOI) to PIC



# Logical Model of PIC Behavior (cont.)


- Map to ports 0x20 & 0x21
  - given ADDR bus
  - logic to form A and CS' inputs to PIC



# Logical Model of PIC Behavior (cont.)


**0x20 => 0000 0000 0010 000 0**

**0x21 => 0000 0000 0010 000 1**

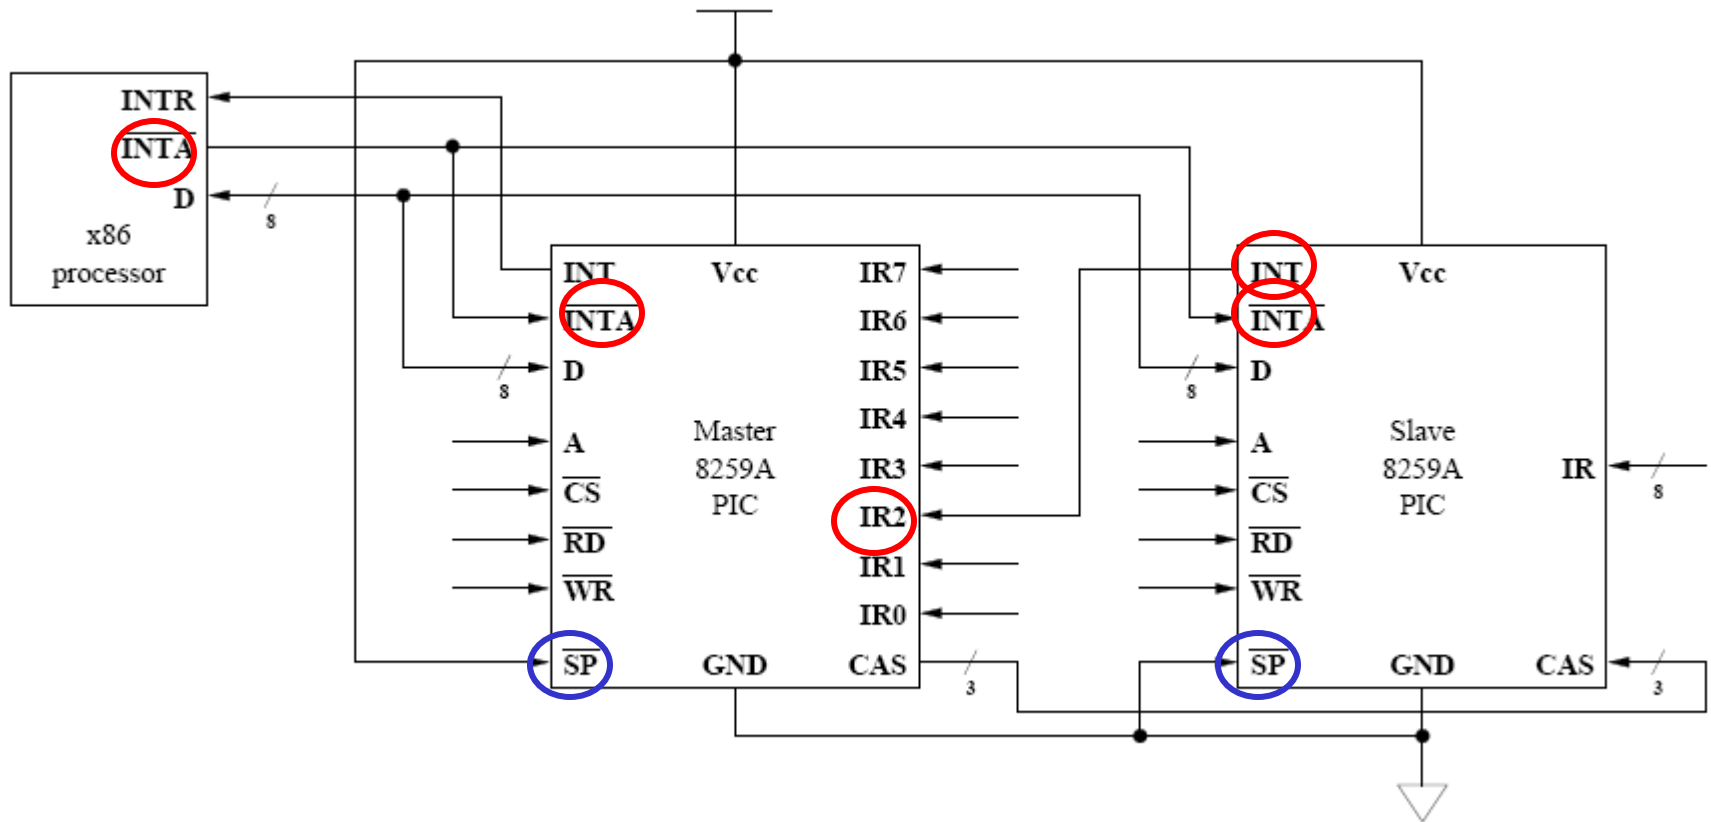


# Logical Model of PIC Behavior (cont.)

- Are eight devices enough? No? What then?
  - hook 2, 3, or 4 PICs to processor?
  - [same problem as before!]
  - design a big PIC?
  - [waste of transistors; not cheap in 8259A era]
  - design several PICs?
  - [waste of humans! (and production costs)]
- Better answer: cascade



# Cascade Configuration of PICs



# Cascading (cont.)

- Previous figure showed x86 configuration of two 8259A's
  - master 8259A mapped to ports 0x20 & 0x21
  - slave 8259A mapped to ports 0xA0 & 0xA1
  - slave connects to IR2 on master
- Question
  - prioritization on 8259A: 0 is high, 7 is low
  - what is prioritization across all 15 pins in x86 layout?
  - (highest) M0...M1...S0...S7...M3...M7 (lowest)

# Cascading (cont.)

- Which PIC should write interrupt vector when processor strobes INTA'?
  - Slave doesn't know which interrupt master reported to CPU
  - Master doesn't know which device connected to slave raised the slave INT line.
- Master tells which slave they should write
  - CAS bus transmits number (3-bits)
  - SP for differentiating between master and slave

# Cascading (cont.)

- In Linux (initialization code to be seen shortly)
  - master IR's mapped to vector #'s 0x20 – 0x27
  - slave IR's mapped to vector #'s 0x28 – 0x2F
  - remember the IDT?

# Interrupt Descriptor Table

|                                      |      |  |                                       |
|--------------------------------------|------|--|---------------------------------------|
| 0x00–0x1F<br><br>defined<br>by Intel | 0x00 | division error                                       |                                       |
|                                      | ⋮    |  |                                       |
|                                      | 0x02 | NMI (non-maskable interrupt)                         |                                       |
|                                      | 0x03 | breakpoint (used by KGDB)                            |                                       |
|                                      | 0x04 | overflow   |                                       |
|                                      | ⋮    |  |                                       |
|                                      | 0x0B | segment not present                                  |                                       |
|                                      | 0x0C | stack segment fault                                  |                                       |
|                                      | 0x0D | general protection fault                             |                                       |
|                                      | 0x0E | page fault   |                                       |
|                                      | ⋮    |  |                                       |
| 0x20–0x27<br><br>master<br>8259 PIC  | 0x20 | IRQ0 — timer chip                                    | example<br>of<br>possible<br>settings |
|                                      | 0x21 | IRQ1 — keyboard                                      |                                       |
|                                      | 0x22 | IRQ2 — (cascade to slave)                            |                                       |
|                                      | 0x23 | IRQ3   |                                       |
|                                      | 0x24 | IRQ4 — serial port (KGDB)                            |                                       |
|                                      | 0x25 | IRQ5   |                                       |
|                                      | 0x26 | IRQ6   |                                       |
|                                      | 0x27 | IRQ7   |                                       |
| 0x28–0x2F<br><br>slave<br>8259 PIC   | 0x28 | IRQ8 — real time clock                               |                                       |
|                                      | 0x29 | IRQ9   |                                       |
|                                      | 0x2A | IRQ10  |                                       |
|                                      | 0x2B | IRQ11 — eth0 (network)                               |                                       |
|                                      | 0x2C | IRQ12 — PS/2 mouse                                   |                                       |
|                                      | 0x2D | IRQ13  |                                       |
|                                      | 0x2E | IRQ14 — ide0 (hard drive)                            |                                       |
|                                      | 0x2F | IRQ15  |                                       |
| 0x30–0x7F                            | ⋮    | APIC vectors available to device drivers             |                                       |
| 0x80                                 | 0x80 | system call vector (INT 0x80)                        |                                       |
| 0x81–0xEE                            | ⋮    | more APIC vectors available to device drivers        |                                       |
| 0xEF                                 | 0xEF | local APIC timer                                     |                                       |
| 0xF0–0xFF                            | ⋮    | symmetric multiprocessor (SMP) communication vectors |                                       |



# Linux 8259A Initialization

```
void init_8259A(int auto_eoi)
{
    unsigned long flags;

    i8259A_auto_eoi = auto_eoi;

    spin_lock_irqsave(&i8259A_lock, flags);

    outb(0xff, 0x21);      /* mask all of 8259A-1 */
    outb(0xff, 0xA1);      /* mask all of 8259A-2 */

    /*
     * outb_p - this has to work on a wide range of PC hardware.
     */
    outb_p(0x11, 0x20);    /* ICW1: select 8259A-1 init */
    outb_p(0x20 + 0, 0x21); /* ICW2: 8259A-1 IR0-7 mapped to 0x20-0x27 */
    outb_p(0x04, 0x21);    /* 8259A-1 (the master) has a slave on IR2 */
    if (auto_eoi)
        outb_p(0x03, 0x21); /* master does Auto EOI */
    else
        outb_p(0x01, 0x21); /* master expects normal EOI */

    outb_p(0x11, 0xA0);    /* ICW1: select 8259A-2 init */
    outb_p(0x20 + 8, 0xA1); /* ICW2: 8259A-2 IR0-7 mapped to 0x28-0x2f */
    outb_p(0x02, 0xA1);    /* 8259A-2 is a slave on master's IR2 */
    outb_p(0x01, 0xA1);    /* (slave's support for AEOI in flat mode
                           is to be investigated) */

    if (auto_eoi)
        /*
         * in AEOI mode we just have to mask the interrupt
         * when acking.
         */
        i8259A_irq_type.ack = disable_8259A_irq;
    else
        i8259A_irq_type.ack = mask_and_ack_8259A;

    udelay(100);          /* wait for 8259A to initialize */

    outb(cached_21, 0x21); /* restore master IRQ mask */
    outb(cached_A1, 0xA1); /* restore slave IRQ mask */

    spin_unlock_irqrestore(&i8259A_lock, flags);
}
```

# Comments on Linux' 8259A Initialization Code

- What is the `auto_eoi` parameter?
  - always = 0, not used in Linux
- Four initialization control words to set up the master 8259A
- Four initialization control words to set up the slave 8259A
- ICW1            0        start init, edge-triggered inputs, cascade mode, 4 ICWs
- ICW2            1        high bits of vector #
- ICW3            1        master: bit vector of slaves; slave: input pin on master
- ICW4            1        ISA=x86, normal/auto EOI

# Comments on Linux' 8259A Initialization Code (cont.)

- What does the “\_p” mean on the “outb” macros?
  - add PAUSE instruction after OUTB; “REP NOP” prior to P4
  - delay needed for old devices that cannot handle processor's output rate
- Critical section spans the whole function; why?
  - avoid other 8259A interactions during initialization sequence
  - (device protocol requires that four words be sent in order)
- Why use \_irqsave for critical section?
  - this code called from other interrupt initialization routines
  - which may or may not have cleared IF on processor