

The Calling Convention (1)

- What is a calling convention?
 - generally: rules for subroutine interface structure
 - specifically
 - how information is passed into subroutine
 - how information is returned to caller
 - who owns registers
 - often specified by vendor so that different compilers' code can work together (it's a CONVENTION)
- Parameters for subroutines
 - pushed onto stack
 - from right to left in C
 - order can be language-dependent

The Calling Convention (2)

- Subroutine return values
 - EAX for up to 32 bits
 - EDX:EAX for up to 64 bits
 - floating-point not discussed
- Register ownership
 - return values can be clobbered by subroutine: EAX and EDX
 - caller-saved: subroutine free to clobber; caller must preserve
 - ECX
 - EFLAGS
 - callee-saved: subroutine must preserve value passed in
 - stack structure: ESP and EBP
 - other registers: EBX, ESI, and EDI

Stack Frames in x86 (1)

- The call sequence
 0. save caller-saved registers (if desired)
 1. **push** arguments onto stack
 2. make the call
 3. **pop** arguments off the stack
 4. restore caller-saved registers

Stack Frames in x86 (2)


- The **callee sequence** (creates the stack frame)
 0. save old base pointer and get new one
 1. save callee-saved registers (always)
 2. make space for local variables
 3. do the function body
 4. tear down stack frame (locals)
 5. restore callee-saved registers
 6. load old base pointer
 7. return

Stack Frames in x86 (3)

- Example of caller code (no caller-saved registers considered)

```
int func (int A, int B, int C);
```

```
func (100, 200, 300);
```



```
PUSHL $300  
PUSHL $200  
PUSHL $100  
CALL func  
ADDL $12,%ESP  
# result in EAX
```

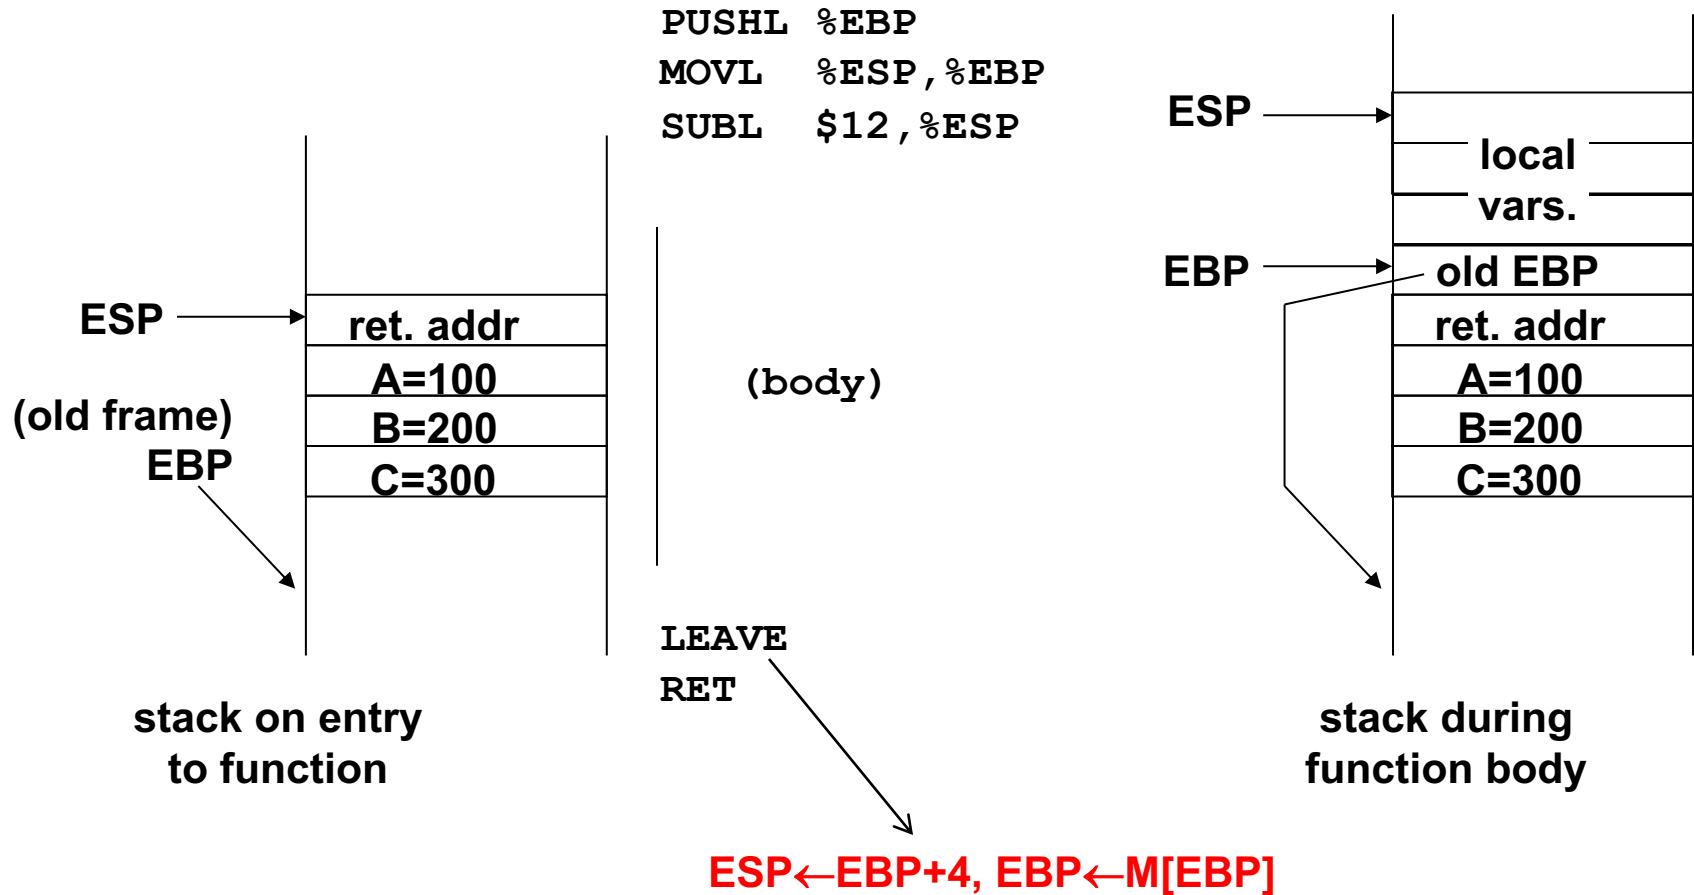
Stack Frames in x86 (4)

- Example of subroutine code and stack frame creation and teardown

```
int func (int A, int B, int C)
{
    /* 12 bytes of local variables */
    ...
}

call func (100, 200, 300);
```

Stack Frames in x86 (4)



Subroutine Example Code

- Earlier assumptions
 - some values start in registers (array pointer in EBX, length in ECX)
 - could specify output regs (min. age in EDX, max. age in EDI)

As a C function, we could write...

```
void find_min_max (person* group, long n_people,  
                  min_max* mm)
```

array of


char* name
long age

long min
long max



```
void find_min_max (person* group, long n_people, min_max* mm)
```



Subroutine Example Code (cont.)



Subroutine Example Code (cont.)



Fun



```
void foo(char *str) {  
    char buffer[256];  
    strcpy(buffer, str);  
}  
  
void main() {  
    char buf[256];  
    memset(buf, 'A', 255);  
    buf[255] = '\\x00';  
    foo(buf);  
}
```


Fun



```
void foo(char *str) {  
    char buffer[16];  
    strcpy(buffer, str);  
}  
  
void main() {  
    char buf[256];  
    memset(buf, 'A', 255);  
    buf[255] = '\\x00';  
    foo(buf);  
}
```


Fun

```
void foo(char *str) {  
    char buffer[16];  
    strcpy(buffer, str);  
}  
  
void main() {  
    char buf[256];  
    memset(buf, 'A', 255);  
    buf[255] = '\x00';  
    foo(buf);  
}
```



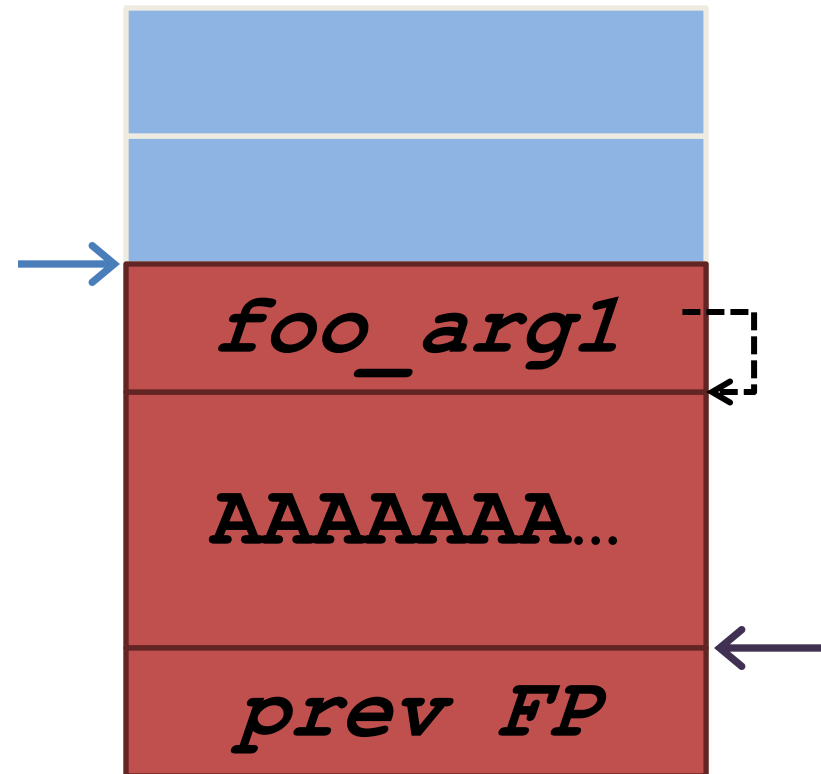
Fun

```
void foo(char *str) {  
    char buffer[16];  
    strcpy(buffer, str);  
}  
  
void main() {  
    char buf[256];  
    memset(buf, 'A', 255);  
    buf[255] = '\\x00';  
    foo(buf);  
}
```




Fun

```
void foo(char *str) {  
    char buffer[16];  
    strcpy(buffer, str);  
}  
  
void main() {  
    char buf[256];  
    memset(buf, 'A', 255);  
    buf[255] = '\x00';  
    foo(buf);  
}
```



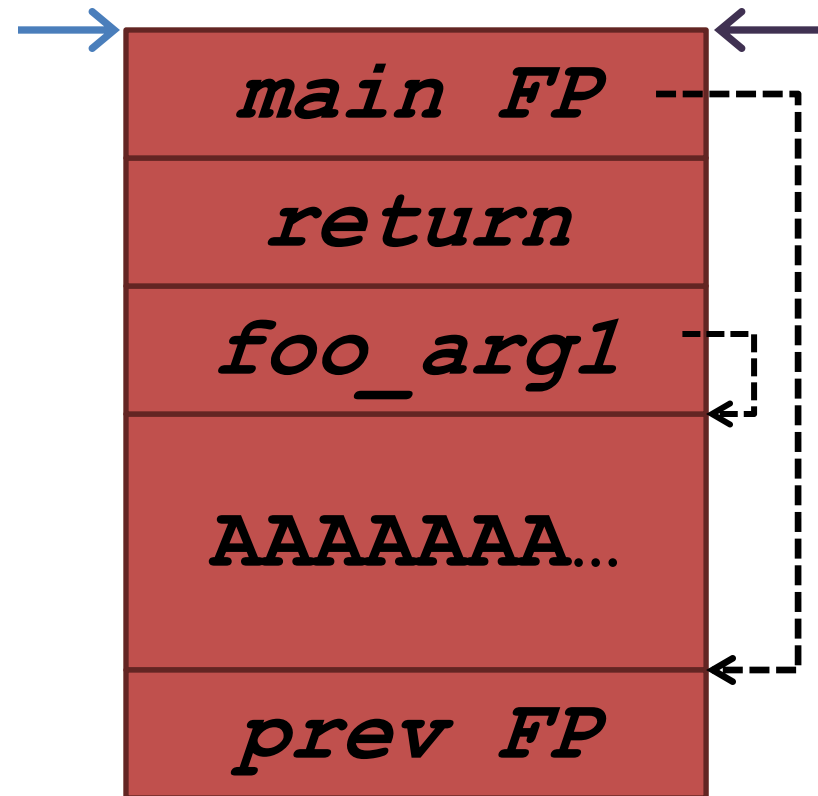
Fun

```
void foo(char *str) {  
    char buffer[16];  
    strcpy(buffer, str);  
}  
  
void main() {  
    char buf[256];  
    memset(buf, 'A', 255);  
    buf[255] = '\x00';  
    foo(buf);  
}
```



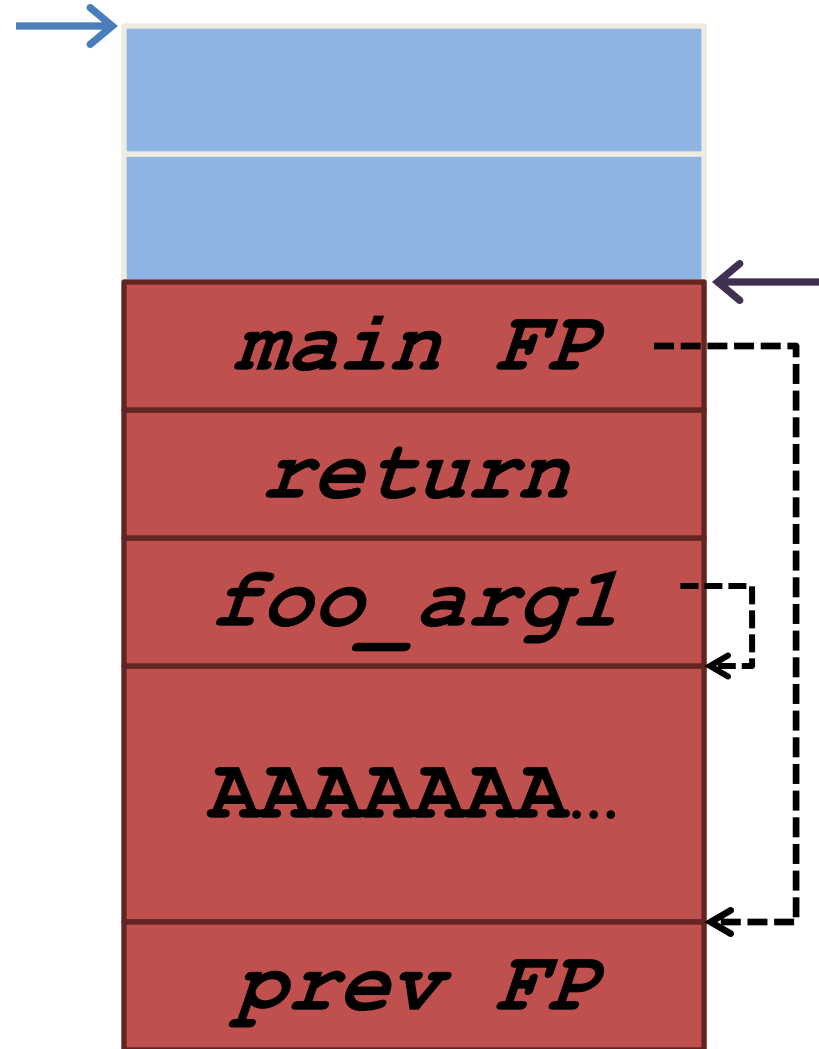
Fun

```
void foo(char *str) {  
    char buffer[16];  
    strcpy(buffer, str);  
}  
  
void main() {  
    char buf[256];  
    memset(buf, 'A', 255);  
    buf[255] = '\x00';  
    foo(buf);  
}
```




Fun

```
void foo(char *str) {  
    char buffer[16];  
    strcpy(buffer, str);  
}  
  
void main() {  
    char buf[256];  
    memset(buf, 'A', 255);  
    buf[255] = '\x00';  
    foo(buf);  
}
```



Fun


```
void foo(char *str) {  
    char buffer[16];  
    strcpy(buffer, str);  
}  
  
void main() {  
    char buf[256];  
    memset(buf, 'A', 255);  
    buf[255] = '\\x00';  
    foo(buf);  
}
```



Fun

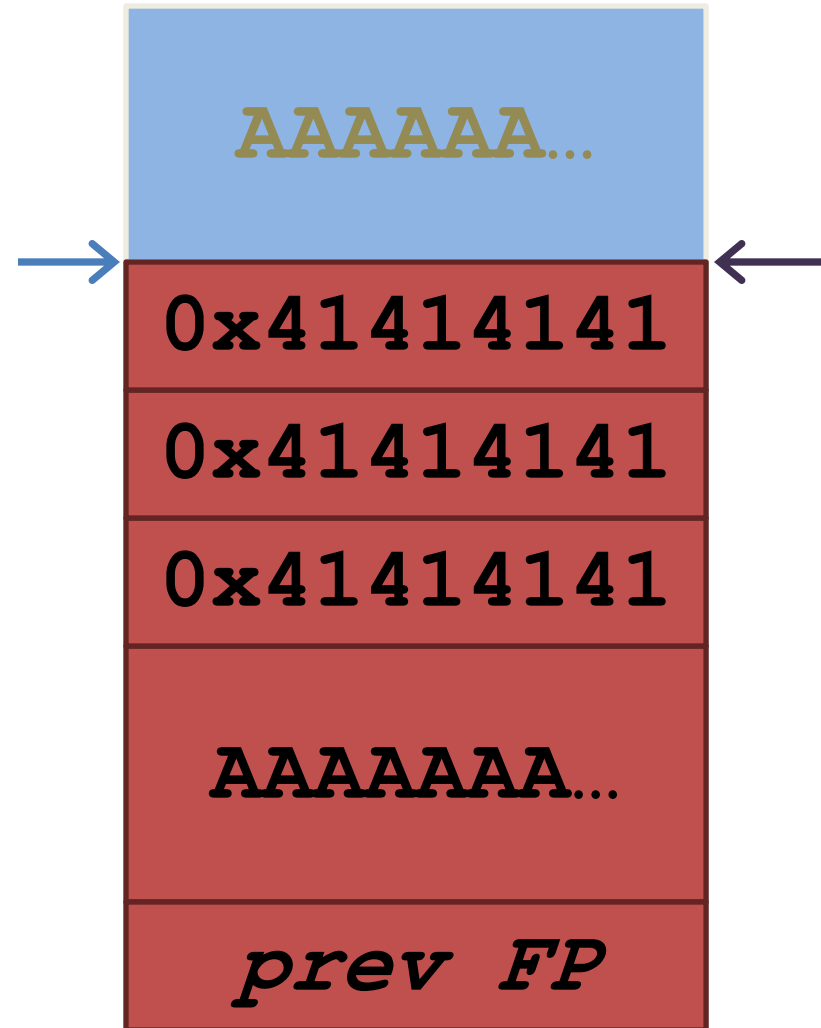
```
void foo(char *str) {  
    char buffer[16];  
    strcpy(buffer, str);  
}  
  
void bar(char *str) {  
    char buf[256];  
    memset(buf, 'A', 255);  
    buf[255] = '\x00';  
    foo(buf);  
}
```

```
mov %ebp, %esp  
pop %ebp  
ret
```



Fun

```
void foo(char *str) {  
    char buffer[16];  
    strcpy(buffer, str);  
    mov %ebp, %esp  
    pop %ebp  
    ret  
}  
void main() {  
    char buf[256];  
    memset(buf, 'A', 255);  
    buf[255] = '\x00';  
    foo(buf);  
}
```




Buffer overflow example

```
void foo(char *str) {  
    char buffer[16];  
    strcpy(buffer, str);  
}  
  
    mov %ebp, %esp  
    pop %ebp  
    ret  
  
void bar() {  
    char buf[256];  
    memset(buf, 'A', 255);  
    buf[255] = '\\x00';  
    foo(buf);  
}
```



Fun

```
void foo(char *str) {  
    char buffer[16];  
    strncpy(buffer, str, 16);  
}  
    mov %ebp, %esp  
    pop %ebp  
    ret  
void bar() {  
    char buf[256];  
    memset(buf, 'A', 255);  
    buf[255] = '\x00';  
    foo(buf);  
}
```




Fun

`%eip = 0x41414141`

???

? ←



Device I/O



- How does a processor communicate with devices?
- Two possibilities
 - independent I/O — use special instructions and a separate I/O port address space
 - memory-mapped I/O — use loads/stores and dedicate part of the memory address space to I/O
- x86 originally used only independent I/O
 - but when used in PC, needed a good interface to video memory
 - solution? put card on the bus, claim memory addresses!
 - now uses both, although ports are somewhat deprecated

- I/O instructions have not evolved since 8086

- 16-bit port space
 - byte addressable
 - little-endian (looks like memory)
- instructions

IN port, dest.reg

OUT src.reg, port


- the register operands are NOT general-purpose registers
 - all data to/from AL, AX, or EAX
 - port is either an 8-bit immediate (not 16!) or DX

Role of System Software (1)

- System software serves three purposes
 - virtualization
 - protection
 - abstraction (particularly hiding asynchrony)
- virtualization:
 - the illusion of multiple/practically unlimited resources
- protection:
 - reduce/eliminate the chance of accidental and/or malicious destruction of data/results by another program


Role of System Software (2)

- abstraction:
 - hide fundamentally asynchronous nature of processor/device interaction
 - provide simpler and more powerful interfaces (integrated w/protection)




System Calls, Interrupts, and Exceptions (1)

- recall that subroutines allow a programmer to encapsulate common operations
- the operating system
 - want to provide an interface including common operations
 - BUT don't want to re-link programs
 - NOR rely on everyone having exactly the same OS version



System Calls, Interrupts, and Exceptions (2)

- solution
 - add a level of indirection!
- with indirection
 - to rewrite OS, just change the table
 - application code does not change



System Calls, Interrupts, and Exceptions (3)

- in LC-3, we used the TRAP instruction; in x86, it's the INT instruction:

INT 8-bit imm. # (PUSH EIP), EIP \leftarrow table[imm8]

- the RTL is actually a little more complicated, as you'll see later in course
- called a trap (after instruction, or trap door through protection boundary)
- also called a system call (for operating system)

System Calls, Interrupts, and Exceptions (4)

- vector tables/jump tables
 - i.e., tables of function pointers
 - convenient abstraction for many procedure-like activities
- Question:
 - What happens if software does something wrong, e.g.,
 - accesses a non-existent memory location?
 - issues an illegal/undefined instruction? divides by 0?
- What do we do to handle problems?
 - state machine that you *design* for processor may have don't cares
 - state machine that you *build* will do something (may be unknown)
 - so just let it run! (e.g., 6502 did so... and programmers used!)

System Calls, Interrupts, and Exceptions (5)


- a better solution: **exceptions!**
 - processor maps each problem to a vector #
 - calls procedure in vector table by #
- Where else might we use vector tables?
- Consider processor interactions with devices
 - a disk access takes about 10 milliseconds
 - new machines in lab: 10 ms = 32 million cycles
- should processor sit around asking, “Are my data here yet?”

System Calls, Interrupts, and Exceptions (6)

- analogous to posting a letter to a friend in Europe
 - and checking your mailbox every minute for a reply
 - instead, have your mail carrier ring your doorbell when it arrives
 - **in a processor, we call that an interrupt**
- How can we use a vector table for interrupts?
 - each device has a vector #
 - call corresponding procedure in vector table when device generates
 - x86 ISA
 - uses one table for all three kinds
 - called the Interrupt Descriptor Table (IDT)

Processor Support for Interrupts (1)

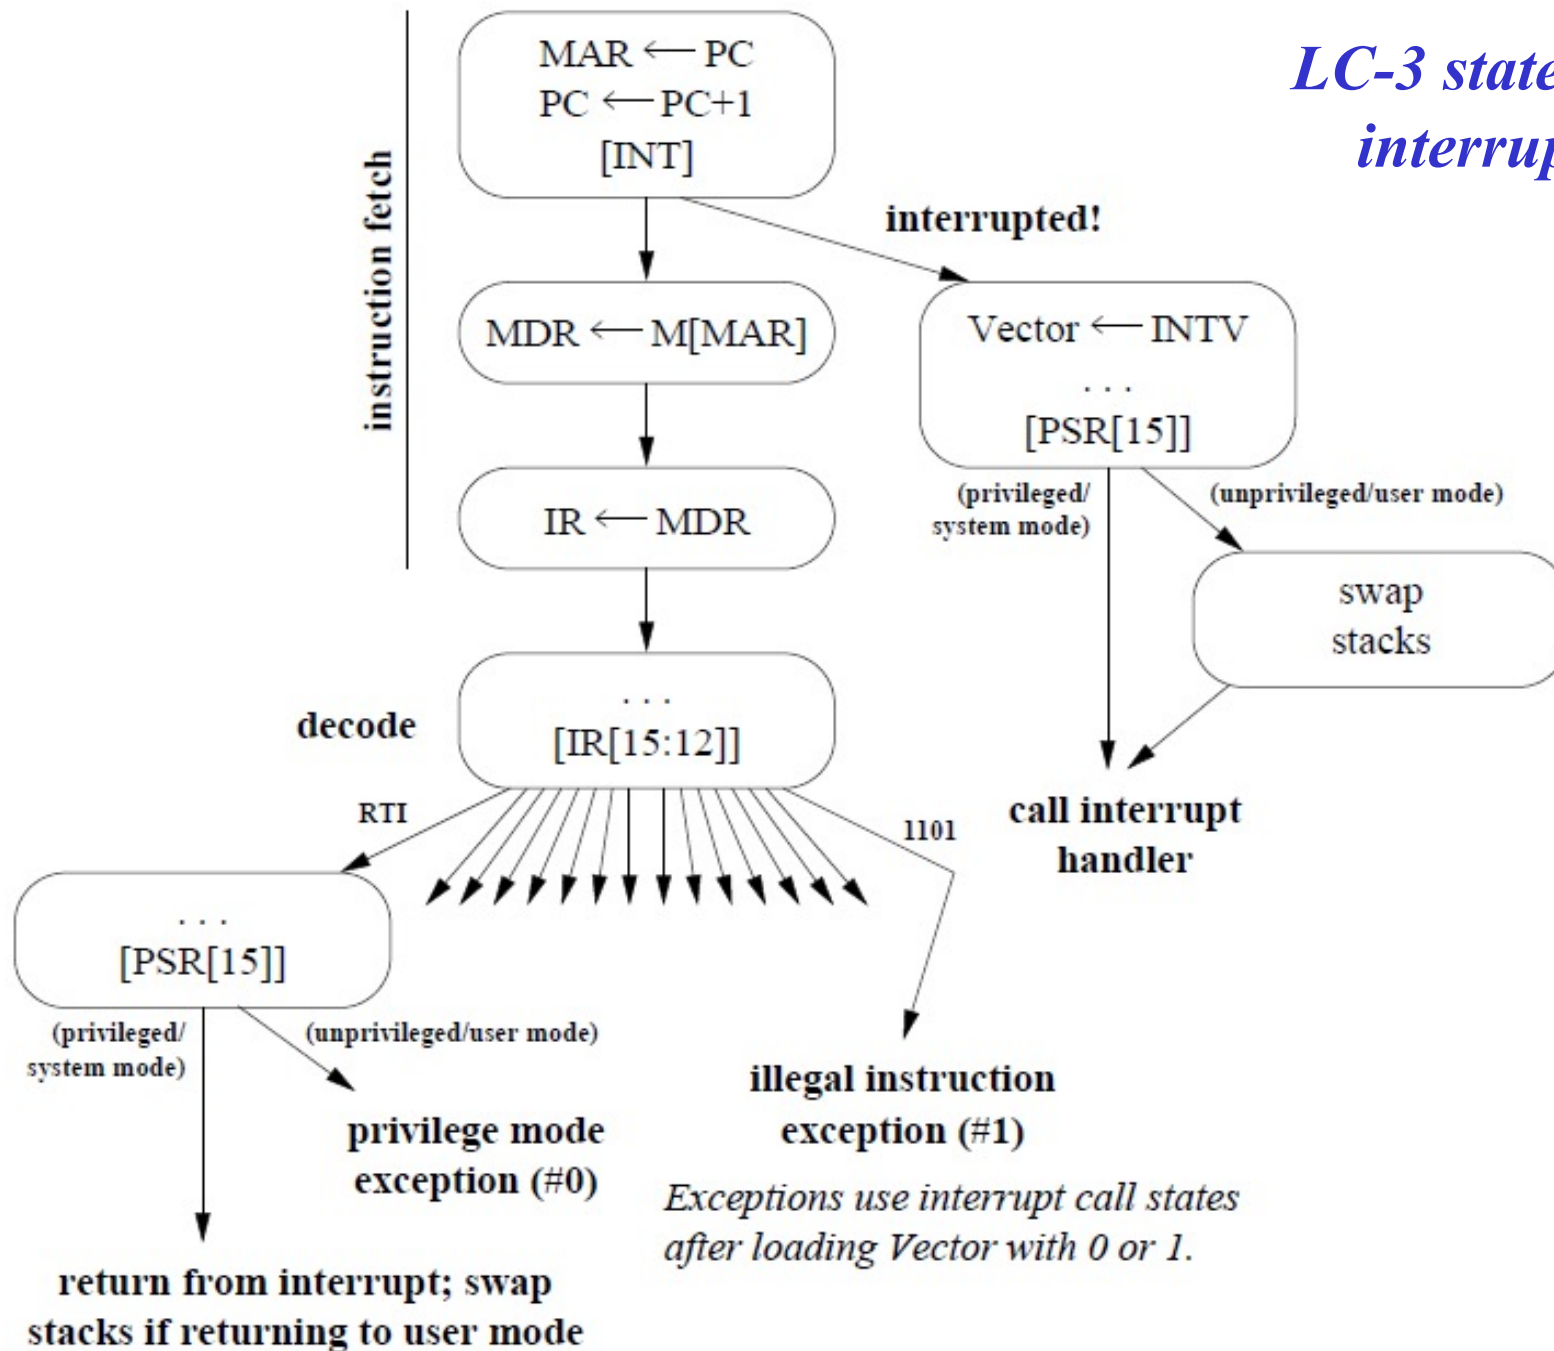
- How does a processor support interrupts?
- Logically...




N=8 on x86 (and LC-3)

- How should we change a processor's state machine to incorporate interrupts?

LC-3 state machine interrupt support



Processor Support for Interrupts (2)



- x86 allows software to block interrupts with a status flag (in EFLAGS)
- normal interrupts occur only when the interrupt enable flag (IF) is set
- some interrupts are too important
 - e.g., memory errors, power warnings, etc.
 - these are NOT maskable, and use a separate input to processor
 - called non-maskable interrupts (NMI)

Interrupt Descriptor Table

- as mentioned earlier
 - x86 uses a single vector table
 - the Interrupt Descriptor Table (IDT)
 - hold vectors for interrupts, exceptions, and system calls
- note that this picture is partly OS-specific
 - the exception vector numbers are specified by Intel – Why?
 - A: generated directly by processor's state machine
 - programmable interrupt controller (PIC) will be discussed later;
 - range of vectors generated is programmable, and is shown for Linux 2.4
 - note that a single entry is used for all system calls in Linux

Interrupt Descriptor Table

0x00–0x1F defined by Intel	0x00	division error	
	⋮		
	0x02	NMI (non-maskable interrupt)	
	0x03	breakpoint (used by KGDB)	
	0x04	overflow	
	⋮		
	0x0B	segment not present	
	0x0C	stack segment fault	
	0x0D	general protection fault	
	0x0E	page fault	
	⋮		
0x20–0x27 master 8259 PIC	0x20	IRQ0 — timer chip	example of possible settings
	0x21	IRQ1 — keyboard	
	0x22	IRQ2 — (cascade to slave)	
	0x23	IRQ3	
	0x24	IRQ4 — serial port (KGDB)	
	0x25	IRQ5	
	0x26	IRQ6	
	0x27	IRQ7	
0x28–0x2F slave 8259 PIC	0x28	IRQ8 — real time clock	
	0x29	IRQ9	
	0x2A	IRQ10	
	0x2B	IRQ11 — eth0 (network)	
	0x2C	IRQ12 — PS/2 mouse	
	0x2D	IRQ13	
	0x2E	IRQ14 — ide0 (hard drive)	
	0x2F	IRQ15	
0x30–0x7F	⋮	APIC vectors available to device drivers	
0x80	0x80	system call vector (INT 0x80)	
0x81–0xEE	⋮	more APIC vectors available to device drivers	
0xEF	0xEF	local APIC timer	
0xF0–0xFF	⋮	symmetric multiprocessor (SMP) communication vectors	

Device I/O



- How does a processor communicate with devices?
- Two possibilities
 - independent I/O — use special instructions and a separate I/O port address space
 - memory-mapped I/O — use loads/stores and dedicate part of the memory address space to I/O
- x86 originally used only independent I/O
 - but when used in PC, needed a good interface to video memory
 - solution? put card on the bus, claim memory addresses!
 - now uses both, although ports are somewhat deprecated

- I/O instructions have not evolved since 8086

- 16-bit port space
 - byte addressable
 - little-endian (looks like memory)
- instructions

IN port, dest.reg

OUT src.reg, port

- the register operands are NOT general-purpose registers
 - all data to/from AL, AX, or EAX
 - port is either an 8-bit immediate (not 16!) or DX